Question 1 (marks = 50)

Shingling and Locality Sensitive Hashing

Problem Statement:

Given:

1. A number of paragraphs from two different books on two different topics.
2. The paragraphs are jumbled up and in no particular order.
3. The paragraphs are of varying length.

Input:

Input will be a text file ***Data.txt*** containing two columns, ***Para No*** (indicates the serial number of the paragraph) and ***Para****.*

To Do:

1. The paragraphs belonging to each book need to be separated based on their similarity.
2. Use Shingling with shingle size **K** = 5.
3. Cluster the similar paragraphs together to reconstruct the book, using k-means algorithm, where **k** = 2.

Output:

The output file produced by your code should be a text file containing the *Para No*s belonging to each book in separate lines. Each Para No belonging to a particular book should be separated by a comma.

Book 1: 1,4,6…
Book 2: 2,3,5….

Question 2 (marks = 50)

Problem Statement:

For the most similar 5 candidate pairs, of each book (set of paragraphs derived above), give the textual overlap regions:

1. The k-shingles that match and
2. Their position indices. Indicate their position in the paragraph, as n^{th} Shingle. If the shingle is present more than once in the paragraph, indicate the first position.

Output:

The output file produced by your code should be a text file containing the following columns

<Para No **m** – Para No **n**><Shingle**1**><Index-**m** – Index-**n**>
<Para No **m** – Para No **n**><Shingle**2**><Index-**m** – Index-**n**>

<Para No **i** – Para No **j**><Shingle**3**><Index-**i** – Index-**j**>
<Para No **i** – Para No **j**><Shingle**4**><Index-**i** – Index-**j**>….